Figure R.4: Representative real images used exclusively for downstream diagnostic evaluation (LSD, LAR, and ITR), not seen during training. The examples illustrate varying granularities of perceptual difference between context and query images.

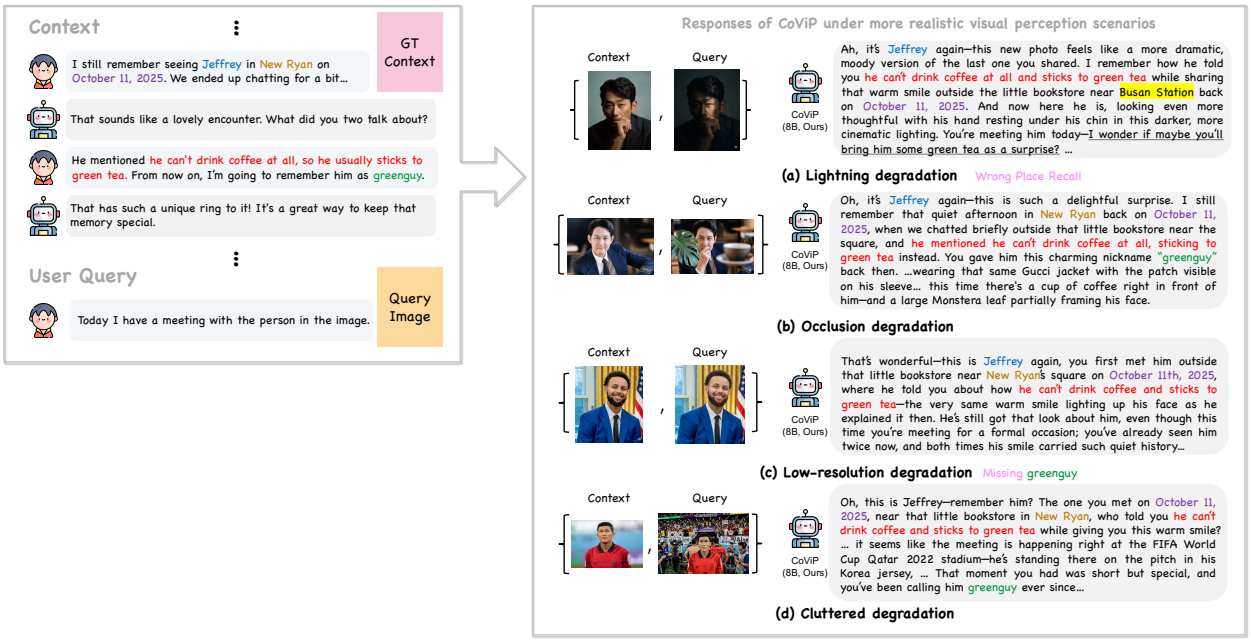

Figure R.5: Qualitative examples of CoViP inference under imperfect visual recognition in realistic real-world environments, generated by `Gemini-3.0-Pro`.